

Lab: LeNet in TensorFlow

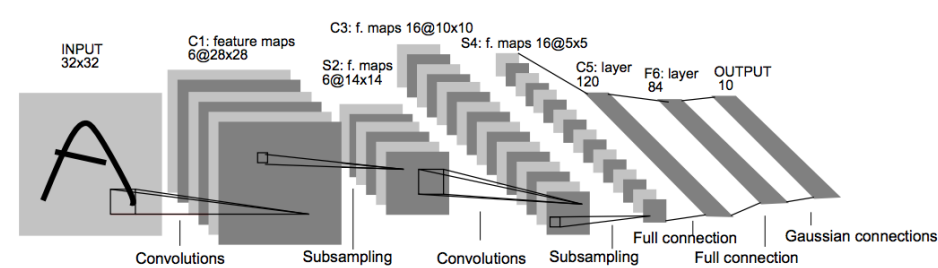


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeNet. Source: Yann Lecun.

You're now going to put together everything you've learned and implement the [LeNet](#) architecture using TensorFlow.

When you get to your next project, remember that LeNet can be a great starting point for your network architecture!

Instructions:

1. Set up your development environment with the [CarND Starter Kit](#)
2. `git clone https://github.com/udacity/CarND-LeNet-Lab.git`
3. `cd CarND-LeNet-Lab`
4. `jupyter notebook`
5. Finish off the architecture implementation in the [LeNet](#) function. That's the only piece that's missing.

Preprocessing

An MNIST image is initially 784 features (1D). If the data is not normalized from [0, 255] to [0, 1], normalize it. We reshape this to `(28, 28, 1)` (3D), and pad the image with 0s such that the height and width are 32 (centers digit further). Thus, the input shape going into the first convolutional layer is `32x32x1`.

Specs

**Convolution layer 1.** The output shape should be `28x28x6`.

**Activation 1.** Your choice of activation function.

**Pooling layer 1.** The output shape should be `14x14x6`.

**Convolution layer 2.** The output shape should be `10x10x16`.

**Activation 2.** Your choice of activation function.

**Pooling layer 2.** The output shape should be `5x5x16`.

**Flatten layer.** Flatten the output shape of the final pooling layer such that it's 1D instead of 3D. The easiest way to do is by using `tf.contrib.layers.flatten`, which is already imported for you.

**Fully connected layer 1.** This should have **120 outputs**.

**Activation 3.** Your choice of activation function.

**Fully connected layer 2.** This should have **84 outputs**.

**Activation 4.** Your choice of activation function.

**Fully connected layer 3.** This should have **10 outputs**.

You'll return the result of the final fully connected layer from the [LeNet](#) function.

If implemented correctly you should see output similar to the following:

```
EPOCH 1 ...
Validation loss = 52.809
Validation accuracy = 0.864

EPOCH 2 ...
Validation loss = 24.749
Validation accuracy = 0.915

EPOCH 3 ...
Validation loss = 17.719
Validation accuracy = 0.930

EPOCH 4 ...
Validation loss = 12.188
Validation accuracy = 0.943

EPOCH 5 ...
Validation loss = 8.935
Validation accuracy = 0.954

EPOCH 6 ...
Validation loss = 7.674
Validation accuracy = 0.956

EPOCH 7 ...
Validation loss = 6.822
Validation accuracy = 0.956

EPOCH 8 ...
Validation loss = 5.451
Validation accuracy = 0.961

EPOCH 9 ...
Validation loss = 4.881
Validation accuracy = 0.964

EPOCH 10 ...
Validation loss = 4.623
Validation accuracy = 0.964

Test loss = 4.726
Test accuracy = 0.962
```

Parameters Galore

As an additional fun exercise calculate the total number of parameters used by the network. Note, the convolutional layers use weight sharing!

Supporting Materials

[lenet.py](#)